

---

# Property Graphs

# The Property Graph Data Model

---

- ❑ Born in the database community
  - Meant to be queried and processed
  - **THERE IS NOTHING SUCH A STANDARD!**
- ❑ Two main constructs: nodes and edges
  - Nodes represent entities,
  - Edges relate pairs of nodes, and may represent different types of relationships
- ❑ Nodes and edges might be labeled,
- ❑ and may have a set of properties represented as attributes (key-value pairs)\*\*\*
- ❑ Further assumptions:
  - Edges are directed,
  - Multi-graphs are allowed

\*\*\* *Note: in some definitions (the least) edges are not allowed to have attributes*

# Formal Definition

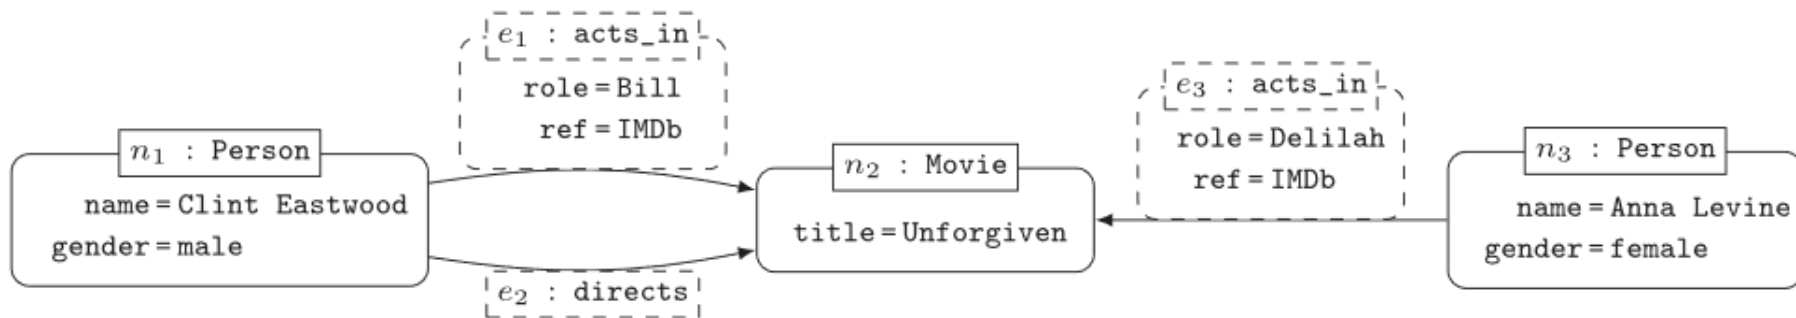
---

*Definition 2.3 (Property graph).* A property graph  $G$  is a tuple  $(V, E, \rho, \lambda, \sigma)$ , where:

- (1)  $V$  is a finite set of *vertices* (or *nodes*).
- (2)  $E$  is a finite set of *edges* such that  $V$  and  $E$  have no elements in common.
- (3)  $\rho : E \rightarrow (V \times V)$  is a total function. Intuitively,  $\rho(e) = (v_1, v_2)$  indicates that  $e$  is a directed edge *from* node  $v_1$  *to* node  $v_2$  in  $G$ .
- (4)  $\lambda : (V \cup E) \rightarrow Lab$  is a total function with  $Lab$  a set of labels. Intuitively, if  $v \in V$  (respectively,  $e \in E$ ) and  $\lambda(v) = \ell$  (respectively,  $\lambda(e) = \ell$ ), then  $\ell$  is the label of node  $v$  (respectively, edge  $e$ ) in  $G$ .
- (5)  $\sigma : (V \cup E) \times Prop \rightarrow Val$  is a partial function with  $Prop$  a finite set of properties and  $Val$  a set of values. Intuitively, if  $v \in V$  (respectively,  $e \in E$ ),  $p \in Prop$  and  $\sigma(v, p) = s$  (respectively,  $\sigma(e, p) = s$ ), then  $s$  is the value of property  $p$  for node  $v$  (respectively, edge  $e$ ) in the property graph  $G$ .

*Extracted from: R. Angles et al. Foundations of Modern Query Languages for Graph Databases*

# Example of Property Graph



Formal definition:

$$V = \{n_1, n_2, n_3\}$$

$$E = \{e_1, e_2, e_3\}$$

$$\rho(e_1) = (n_1, n_2)$$

$$\rho(e_2) = (n_1, n_2)$$

$$\rho(e_3) = (n_3, n_2)$$

$$\lambda(n_1) = \text{Person}$$

$$\lambda(n_2) = \text{Movie}$$

$$\lambda(n_3) = \text{Person}$$

$$\lambda(e_1) = \text{acts\_in}$$

$$\lambda(e_2) = \text{directs}$$

$$\lambda(e_3) = \text{acts\_in}$$

$$\sigma(n_1, \text{name}) = \text{Clint Eastwood}$$

$$\sigma(n_1, \text{gender}) = \text{male}$$

$$\sigma(n_2, \text{title}) = \text{Unforgiven}$$

$$\sigma(n_3, \text{name}) = \text{Anna Levine}$$

$$\sigma(n_3, \text{gender}) = \text{female}$$

$$\sigma(e_1, \text{role}) = \text{Bill}$$

$$\sigma(e_1, \text{ref}) = \text{IMDb}$$

$$\sigma(e_3, \text{role}) = \text{Delilah}$$

$$\sigma(e_3, \text{ref}) = \text{IMDb}$$

# Traversal Navigation

---

- We define the graph traversal pattern as:  
*"the ability to rapidly traverse structures to an arbitrary depth (e.g., tree structures, cyclic structures) and with an arbitrary path description (e.g. friends that work together, roads below a certain congestion threshold)"* [Marko Rodriguez]

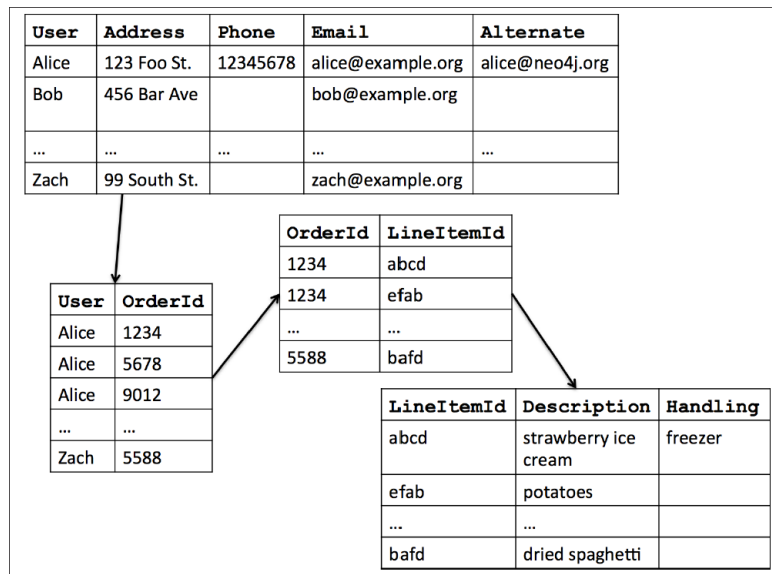
# Traversal Navigation

---

- We define the graph traversal pattern as:  
*"the ability to **rapidly** traverse structures to an **arbitrary depth** (e.g., tree structures, cyclic structures) and with an **arbitrary path description** (e.g. friends that work together, roads below a certain congestion threshold)"* [Marko Rodriguez]
- Totally opposite to set theory (on which relational databases are based on)
  - Sets of elements are operated by means of the relational algebra

# Traversing Data in a RDBMS

- In the relational theory, it is equivalent to joining data (schema level) and select data (based on a value)



```
SELECT *  
FROM user u, user_order uo,  
orders o, items i  
WHERE u.user = uo.user AND  
uo.orderId = o.orderId AND  
i.lineItemId = i.LineItemId  
AND u.user = 'Alice'
```

## Cardinalities:

|User|: 5.000.000

|UserOrder|: 100.000.000

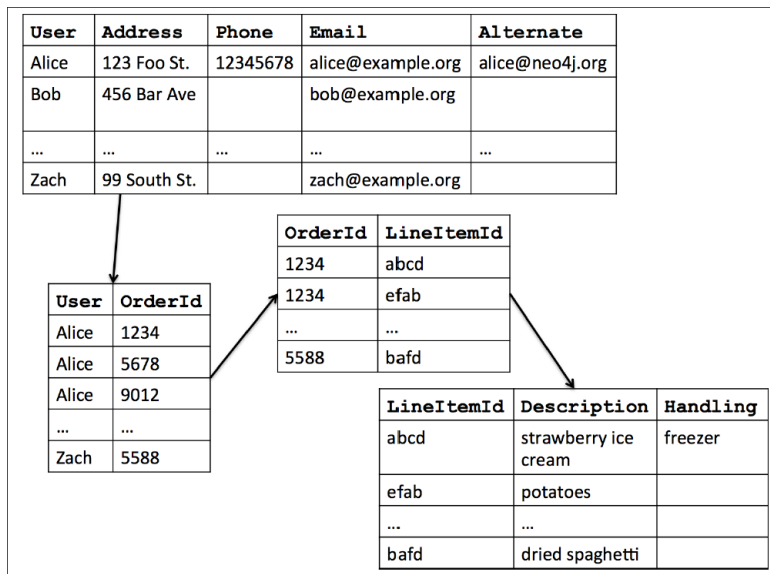
|Orders|: 1.000.000.000

|Item|: 35.000

## Query Cost?!

# Activity

- *Wear your data steward hat and discuss in pairs the database tuning that would guarantee the most efficient access plan for this query*
  - *What join algorithm would you take? Why?*



```
SELECT *  
FROM user u, user_order uo,  
orders o, items i  
WHERE u.user = uo.user AND  
uo.orderId = o.orderId AND  
i.lineItemId = i.LineItemId  
AND u.user = 'Alice'
```

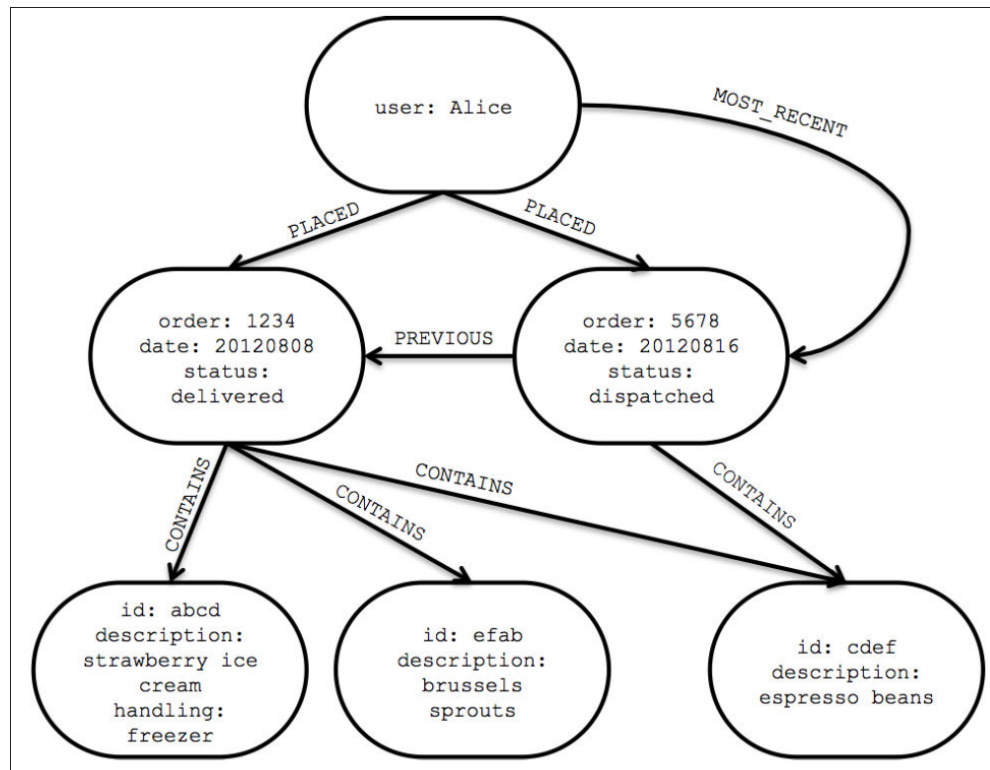
## Cardinalities:

|User|: 5.000.000  
|UserOrder|: 100.000.000  
|Orders|: 1.000.000.000  
|Item|: 35.000

## Query Cost?!



# Traversing Data in a Graph Database



## Cardinalities:

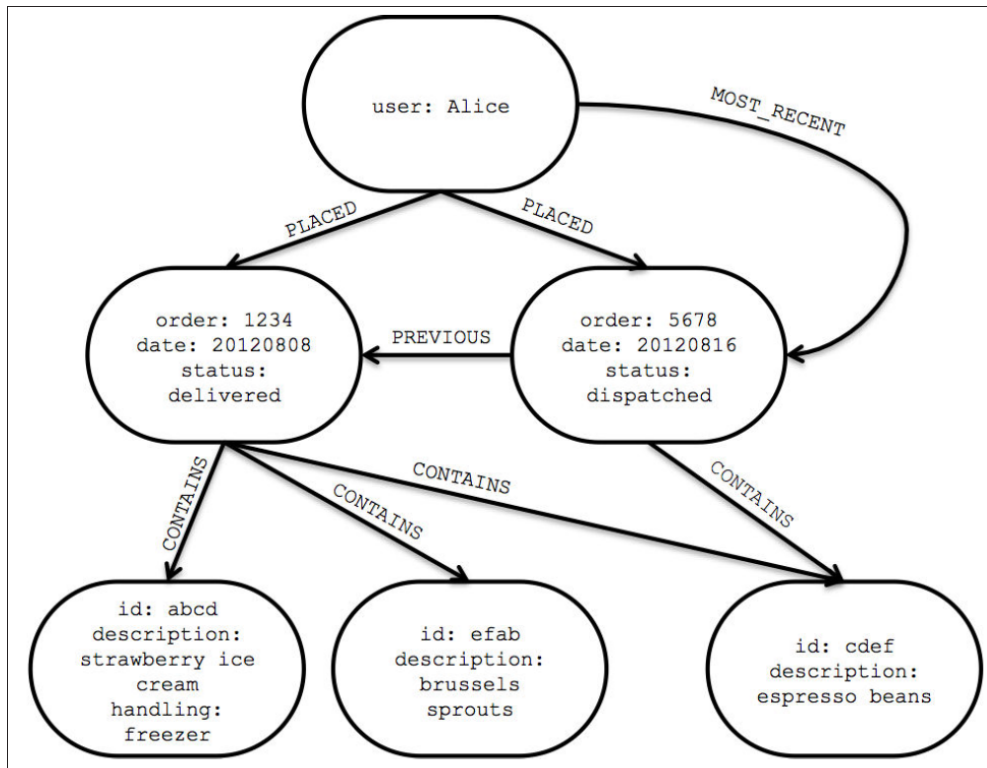
|User|: 5.000.000

|Orders|: 1.000.000.000

|Item|: 35.000

# Activity

- What would be the cost of this query in a graph database?
  - Assume you can find a node in constant time



## Cardinalities:

|User|: 5.000.000

|Orders|: 1.000.000.000

|Item|: 35.000

**Query Cost?!**

# Traversing Property Graphs

---

- Traversing graph data depends on three main variables
  - The size of the graph (i.e., #edges),
  - The topology of the graph,
  - The query topology

---

REFRESHING SOME BASICS ON GRAPHS

# GRAPH OPERATIONS

# Graph Operations

---

## □ Content-based queries

### ■ The value is relevant

- Get a node, get the value of a node / edge attribute, etc.
- A typical case are summarization queries (i.e., aggregations)

## □ Topological queries

### ■ Only the graph topology is considered

### ■ Typically, several business problems (such as fraud detection, trend prediction, product recommendation, network routing or route optimization) are solved using graph algorithms exploring the graph topology

- Computing the betweenness centrality of a node...
  - in a social network, an analyst can detect influential people or groups for targeting a marketing campaign audience.
  - in a telecommunication operator, an analyst may detect central nodes of an antenna network and optimize the routing and load balancing across the infrastructure accordingly

## □ Hybrid approaches

# Topological Queries

---

- Divided in three (four) main categories
  - Adjacency,
  - Reachability,
  - Pattern Matching,
  - [Graph metrics]

# Adjacency Queries

---

- Formalized as node adjacency or edge incidence
  - Node adjacency
  - Edge incidence (node degree, out-degree, in-degree)
  - K-neighbourhood of a node
- Formal definition:

$$\text{Adjacency}(n) = \bar{N}$$

$$n_i \in \bar{N} \iff \exists e_1 \mid \rho(e_1) = (n_i, n) \vee \rho(e_1) = (n, n_i)$$

- Computational cost: linear cost on the number of edges to visit
- Examples:
  - Find all friends of a person
  - Airports with a direct connection
  - Movies watched by a person
  - Products bought by a customer
  - ...

# Reachability Queries

---

## □ Formal definition:

Reachability( $n_{or}, n_{dest}$ ) is **true**  $\iff \exists \text{Walk}(n_{or}, n_{dest})$

$\text{Walk}(n_{or}, n_{dest}) = (e_1 \dots e_m) \mid \exists n_1 \dots n_{m-1}, \rho(e_1) = (n_{or}, n_1), \rho(e_2) = (n_1, n_2) \dots \rho(e_m) = (n_{m-1}, n_{dest})$

## □ Reachability queries **find a walk** between two nodes. They may define additional constraints (e.g., shortest path, fixed-length paths, etc.)

- Fixed-length paths (fixed #edges and nodes)
- Regular simple paths (restrictions as regular expressions)
  - Hybrid if the restriction is in the *content*
- Shortest path
- Non-repeated nodes (path)
- Non-repeated edges (trail)



# Reachability Queries

---

- Computational cost: hard to compute for large graphs
  - Shortest-path (Dijkstra's algorithm):  $O(|V|^2)$ 
    - Smarter implementations based on priority queues yield  $O(|E| * |V| \log |V|)$  complexity
- Examples:
  - Friend-of-a-friend
  - Flight connections
  - Logistics (goods distribution)
  - Items bought in a user orders
  - ...

# Single-Source Shortest-Path

---

## □ Dijkstra's algorithm

### ■ Main idea:

- Optimal substructure: The subpath of any shortest path is itself a shortest path
- Triangle inequality:  $\delta(u,v) \leq \delta(u,x) + \delta(x,v)$ , if  $u,v$  is the shortest path

### ■ Input:

- A weighted graph  $\mathbf{G} = (\mathbf{V}, \mathbf{E})$ ,
- A source vertex  $\mathbf{V}_s \in \mathbf{V}$

### ■ Internal structures:

- $\mathbf{Q}$ : Vertices not yet processed,
- $\mathbf{V-Q}$ : Set of vertices whose shortest paths from the source have already been determined,
- **dist[]**: current estimated shortest paths to each vertex,
- **[prev[]]**: *array of predecessors for each vertex* – traceback

### ■ Output:

- The graph with all the distances from  $\mathbf{V}_s$  to all nodes in  $\mathbf{V} - \mathbf{V}_s$ 
  - The graph representing all the paths from one vertex to all the others must be a **spanning tree** (minimum number of edges)
  - There will be no cycles as a cycle would define more than one path from the selected vertex to at least one other vertex

# Pseudocode

---


```
1  function Dijkstra(Graph, source):
2
3      create vertex set Q
4
5      for each vertex v in Graph:
6          dist[v] ← INFINITY
7          prev[v] ← UNDEFINED
8          add v to Q
9
10     dist[source] ← 0
11
12     while Q is not empty:
13         u ← vertex in Q with min dist[u]
14
15         remove u from Q
16
17         for each neighbor v of u:           // only v that are still in Q
18             alt ← dist[u] + length(u, v)
19             if alt < dist[v]:
20                 dist[v] ← alt
21
22
23     return dist[]
```

# Pseudocode

Animation by Ibmuu

<https://commons.wikimedia.org/w/index.php?curid=6282617>


```
1  function Dijkstra(Graph, source):
2
3      create vertex set Q
4
5      for each vertex v in Graph:
6          dist[v] ← INFINITY
7          prev[v] ← UNDEFINED
8          add v to Q
9      dist[source] ← 0
10
11
12     while Q is not empty:
13         u ← vertex in Q with min dist[u]
14
15         remove u from Q
16
17         for each neighbor v of u:           // only v that are still in Q
18             alt ← dist[u] + length(u, v)
19             if alt < dist[v]:
20                 dist[v] ← alt
21
22
23     return dist[]
```



# Activity

---

- ❑ Modify Dijkstra's algorithm such that it also returns the path between the source and every other node
- ❑ Compute Dijkstra's algorithm (returning the paths) on the following example:



# Label-constrained Reachability

---

## □ Definition:

$G_L^* = \{(s, t) \mid \text{there is a path in } G \text{ from } s \text{ to } t \text{ using only edges with labels in } L\}$

- It is equivalent to determine whether or not there is a path in **G** from  $s$  to  $t$  such that the concatenation of the edge labels along the path forms a string in the language denoted by the regular expression  $(\ell_1 \cup \dots \cup \ell_n)^*$

where:  $L = \{\ell_1, \dots, \ell_n\}$ ,  $\cup$  disjunction and  $*$  the Kleene star

- Typically, the allowed topology and labels involved are expressed as a regular expression
  - In general, the cost of regular label-constrained queries is known to be NP-complete

# Pattern Matching

---

- Formalized as the graph isomorphism problem
  - Input: property graph **G**, and a pattern graph **P**
  - Output: all sub-graphs of **G** that are isomorphic to **P**
- Computational cost: hard to compute, in general, NP-complete
- Examples:
  - Group of cities all of them directly connected by flights
  - People without telephone
  - People who have watched sci-fi movies in the last month
  - ...

# Property Graph Patterns

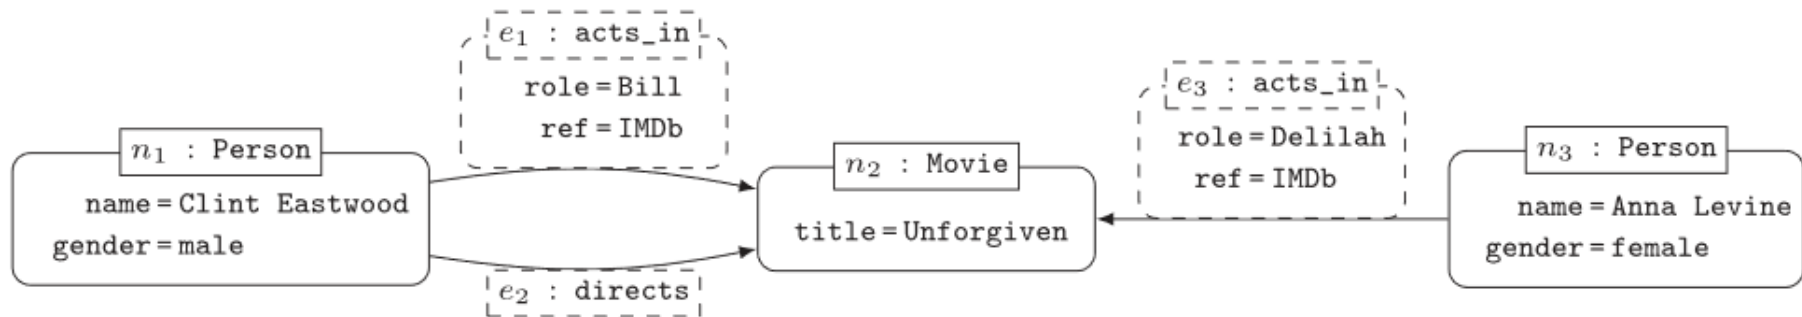
---

- Based on *basic graph patterns* (bgps)
  - Equivalent to conjunctive queries
- A *bgp* for querying property graphs is a property graph where variables can appear in place of any constant (labels / properties)
- A **match** for a *bgp* is a mapping from variables to constants such that when the mapping is applied to the *bgp*, the result is *contained* within the original graph
- The **results** for a *bgp* are then all mappings from variables in the query to constants that comprise a match




# Example of Graph Pattern

## Graph:



## BGP:



# Evaluating Graph Patterns

---

- Evaluating a *bgp*  $Q$  against a graph database  $G$  corresponds to listing all possible matches of  $Q$  with respect to  $G$
- Formally:

*Definition 3.5 (Match).* Given an edge-labelled graph  $G = (V, E)$  and a *bgp*  $Q = (V', E')$ , a *match*  $h$  of  $Q$  in  $G$  is a mapping from  $Const \cup Var$  to  $Const$  such that:

- (1) for each constant  $a \in Const$ , it is the case that  $h(a) = a$ ; that is, the mapping maps constants to themselves; and
- (2) for each edge  $(b, l, c) \in E'$ , it holds that  $(h(b), h(l), h(c)) \in E$ ; this condition imposes that (a) each edge of  $Q$  is mapped to an edge of  $G$ , and (b) the structure of  $Q$  is preserved in its image under  $h$  in  $G$  (that is, when  $h$  is applied to all the terms in  $Q$ , the result is a sub-graph of  $G$ ).

*Extracted from: R. Angles et al. Foundations of Modern Query Languages for Graph Databases*

# Semantics of a Match

---

- **Homomorphism-based semantics:** the previous definition maps to a homomorphism from  $\mathbf{Q}$  to  $\mathbf{G}$ 
  - Multiple variables in  $\mathbf{Q}$  can map to the same term in  $\mathbf{G}$
  - Corresponds to the familiar semantics of select-from-where queries (conjunctive queries) in relational databases
- **Isomorphism-based semantics:** an additional constraint is added, the match function  $h$  must be injective.

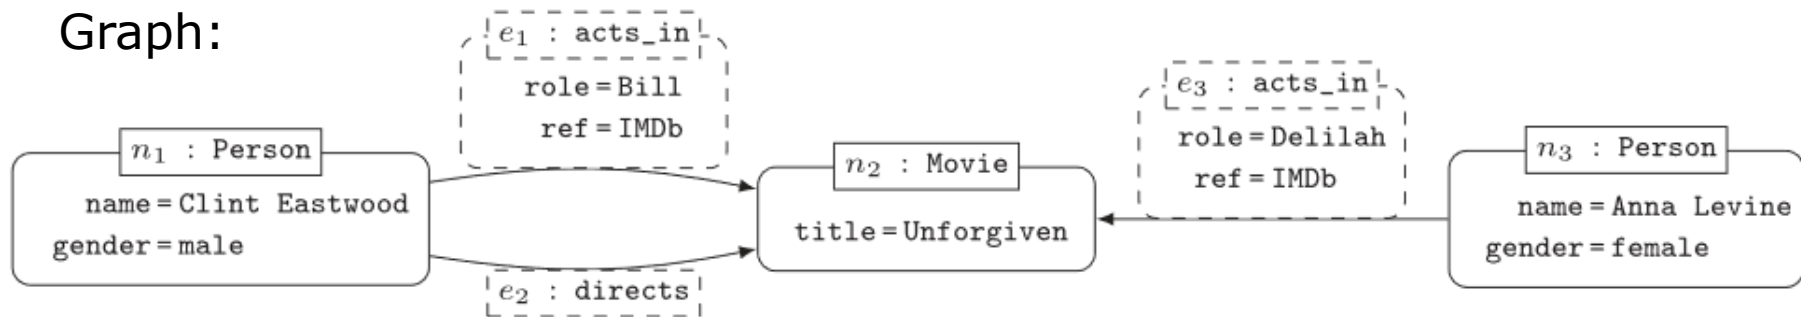
Still, different semantics can be applied:

- Strict isomorphism (no-repeated-anything):  $h$  is injective
- No repeated-node semantics:  $h$  is only injective for nodes
- No repeated-edge semantics:  $h$  is only injective for edges


# Activity

- Objective: Understand the differences between graph matching isomorphism-based and homomorphism-based semantics
  - Given the following graph, bgp and potential results...

Graph:



BGP:



# Activity

- *Objective: Understand the differences between graph matching isomorphism-based and homomorphism-based semantics*
  - *Given the following graph, bgp and potential results...*

Results:

$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$x_6$	$x_7$	$x_8$	$x_9$
$n_2$	$e_2$	directs	$n_1$	Clint Eastwood	$e_3$	acts_in	$n_3$	Anna Levine
$n_2$	$e_3$	acts_in	$n_3$	Anna Levine	$e_2$	directs	$n_1$	Clint Eastwood
$n_2$	$e_1$	acts_in	$n_1$	Clint Eastwood	$e_3$	acts_in	$n_3$	Anna Levine
$n_2$	$e_3$	acts_in	$n_3$	Anna Levine	$e_1$	acts_in	$n_1$	Clint Eastwood
$n_2$	$e_2$	directs	$n_1$	Clint Eastwood	$e_1$	acts_in	$n_1$	Clint Eastwood
$n_2$	$e_1$	acts_in	$n_1$	Clint Eastwood	$e_2$	directs	$n_1$	Clint Eastwood
$n_2$	$e_1$	acts_in	$n_1$	Clint Eastwood	$e_1$	acts_in	$n_1$	Clint Eastwood
$n_2$	$e_2$	directs	$n_1$	Clint Eastwood	$e_2$	directs	$n_1$	Clint Eastwood
$n_2$	$e_3$	acts_in	$n_1$	Anna Levine	$e_3$	acts_in	$n_1$	Anna Levine

**Circle what results would be obtained if applying isomorphism-based or homomorphism-based semantics**

- **Distinguish the three isomorphism semantics presented**

# From Intractable to Tractable Matching

---

- These results apply to property graphs:
  - Graph isomorphism is known to be NP-complete in the worst case
  - Graph homomorphism is also known to be NP-complete in the worst case
  - However, *graph simulation and bi-simulation*, a relaxed form of graph homomorphism, can be computed within polynomial time
    - It might still hard to compute for large graphs
    - New iterative algorithms allow to scale-well

*Fan et al. Graph Pattern Matching: From Intractable to Polynomial Time. VLDB'10*

(<https://www.comp.nus.edu.sg/~vldb2010/proceedings/files/papers/R23.pdf>)

# Graph Metrics

---

- They can be formalized either as adjacency, reachability or pattern matching
  - Thus, the cost depends on how the metric is formalized
- Given their relevance, they are typically provided as built-in functions
- Examples:
  - Graph node order,
  - the min / max degree in the graph,
  - the length of a path,
  - the graph diameter,
  - the graph density,
  - closeness / betweenness of a node,
  - the pageRank of a node,
  - ...

# Graph Metrics

---

- In this course, we will focus on the following algorithms (1<sup>st</sup> lab):
  - Centrality algorithms
    - Page Rank
    - Betweenness
    - Closeness
  - Path finding algorithms
    - Minimum weight spanning tree
    - Single source shortest path
  - Community detection algorithms
    - Triangle counting
    - Louvain
    - (Strongly) connected components



# Topological Queries

## □ Support provided by current GBDs

<i>Graph Database</i>	Adjacency		Reachability			Pattern matching	Summarization
	Node/edge adjacency	k-neighborhood	Fixed-length paths	Regular simple paths	Shortest path		
Allegro	•		•			•	
DEX	•		•	•	•	•	
Filament	•		•			•	
G-Store	•		•	•	•	•	
HyperGraph	•					•	
Infinite	•		•	•	•	•	
Neo4j	•		•	•	•	•	
Sones	•					•	
vertexDB	•		•	•		•	

R. Angles. A Comparison of Current Graph Database Models (as of 2012)

# Implementation of the Operations

---

- Note that the operations presented are conceptual: agnostic of the technology
- The implementation of the ops depends on:
  - The graph database implementation,
  - The operation implementation,
  - The pattern (in case of pattern matching)

# Summary

---

- Property graphs do not have a defined standard, but there is a de-facto standard
  - Nodes / edges may have labels (equivalent to the concept of typing) and / or properties (to the concept of attributes)
- The basic operations on graphs are:
  - Adjacency queries,
  - Reachability queries and
  - Pattern matching
- However, their traditional definition needs to be redefined for property graphs, yielding different computational complexities
  - It basically depends on the graph topology, graph pattern and internal structures the graph database is implemented on